PARAG GUPTA

Carson, CA 90746 | paragguptaus@gmail.com | (732) 558-4884 | github.com/ParagGupta103

EDUCATION

Rutgers University, School of Arts & Sciences | Honors College | Aug 2022 – May 2026
Bachelor of Science in Computer Science (GPA: 3.6/4.0) | New Brunswick, NJ

- Honors: Presidential Scholar, SAS Excellence Award, Dean's List Award
- Relevant Coursework: Data Structures & Algorithms, Massive Data Mining, Data Management for Data Science, Applied Statistical Learning

SUMMARY

Software Development Engineer I candidate with **2+ years of internship experience** building scalable, cloud-native systems in Python, Java, and AWS. Reduced pipeline latency by **30%** by architecting scalable Python and SQL ETL pipelines on GCP, applying data structures, algorithms, and object-oriented design. Builds distributed systems and GenAI solutions, using Claude Code to accelerate development and debugging.

EXPERIENCE

KW International | Aug 2026 – Present
Data & Automation Analyst | Carson, CA

- Automated last-mile shipping quotes by building Python ETL pipelines that integrate **4 carrier APIs** (FedEx, UPS, USPS, DHL), saving **10+ hours per week**.
- Developed a warehouse cost estimation model that projects storage, palletizing, loading, and fuel-inclusive transport costs from 3PL shipment history, achieving **100% accuracy** when given a detailed price breakdown.

Collaborative Solutions | Jan 2026 – May 2026
Data Engineering Extern | New Brunswick, NJ

- Engineered an automated Python spatial ETL pipeline that ingested multi-angle imagery for **500+ rail crossings** via the Google Maps API and converted it into safety metrics.
- Deployed a custom YOLOv5 model and risk-scoring engine that detected clutter with **85% precision**, aggregating detections into a scalable cloud framework for statewide infrastructure analysis.

Tarifflo | May 2025 – Aug 2025
Data Engineering Intern | Provo, UT

- Architected scalable Python and SQL ETL pipelines on GCP for real-time logistics data, reducing pipeline latency by **30%** for downstream BigQuery analytics and Vertex AI models.
- Modeled the Harmonized Tariff Schedule and built customs valuation models for CBP's ACE system, covering all HTS codes with **98% accuracy**.

Newmark Group | Sep 2024 – Dec 2024
Corporate Analyst Program | Rutherford, NJ

- Built modular Python pipelines that converted fragmented financial data into Parquet assets and fed GCP Looker dashboards, replacing manual data entry and saving **15+ hours per week**.

Apollo Tyres | May 2024 – Jul 2024
Software Engineering Intern | New Delhi, India

- Built a GenAI RAG pipeline on AWS Bedrock with Python validation scripts, making **500+ scanned documents** searchable.
- Reduced query latency by **20%** by troubleshooting and refactoring OpenCypher queries in AWS Neptune, a NoSQL graph database.

PROJECTS

MotoLens – Shazam for Cars | *Python, PyTorch, FastAPI, React, PostgreSQL*

- Built a full-stack computer vision platform serving a PyTorch ResNet-50 model through a FastAPI backend, covering **196 car classes** trained on **about 10,000 images**.

Stock Market Data Warehouse | *Apache Airflow, dbt, Snowflake, GCP BigQuery, Python*

- Reduced compute costs by **40%** by orchestrating Airflow pipelines into Snowflake and BigQuery with optimized dbt models and data-quality checks.

SKILLS

Automation Framework: pytest, Selenium, Playwright
Programming Languages: Python, Java, C++, SQL
API Testing Tools: Postman
CI/CD Tools: GitLab, Git, Docker, Kubernetes, Terraform, Airflow
Cloud Platforms: AWS (Bedrock, Neptune), GCP (BigQuery, Vertex AI, Looker)
Databases: SQL (PostgreSQL, BigQuery SQL, Snowflake), NoSQL (MongoDB, AWS Neptune)